

Name: **Serene Plummer**

Assignment: Week 8 Review Lab

Date Due: Saturday, October 11, 2025

Class: Computer Science II - CSCE 1040 Section 501

<https://github.com/serene4444/Estimating-Housing-Prices-using-Support-Vector-Machine>

```
$ python build\M5-Problem.py
C:\Users\Mercu\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\sklearn\svm\_base.py:305: ConvergenceWarning: Solver terminated early (max_iter=1000). Consider pre-processing your data with StandardScaler or MinMaxScaler.
    warnings.warn(
Training SVR model...
Training completed!
Mean squared error:  266.50178022397176
Explained variance score:  -179.7668670020227

Confusion Matrix:
[[170 389]
 [ 96 345]]
precision    recall   f1-score   support
          0       0.64      0.30      0.41      559
          1       0.47      0.78      0.59      441

   accuracy                           0.52      1000
  macro avg       0.55      0.54      0.50      1000
weighted avg       0.56      0.52      0.49      1000

Threshold: 2.0
Actual labels - Class 0: 559, Class 1: 441
Predicted labels - Class 0: 266, Class 1: 734
```

The screenshot shows a Jupyter Notebook environment with the following details:

- Code Editor:** The file `M5-Problem.py` is open, showing Python code for importing packages, loading datasets, and performing a linear SVM regression.
- Output Area:** A figure titled "Confusion matrix" is displayed, showing a 2x2 matrix with values 170, 389, 96, and 345. Below the figure, the text "Mean squared error: 266.50178022397176" and "Explained variance score: -179.7668670020227" is shown.
- Text Output:** The confusion matrix is printed as a table:

precision	recall	f1-score	support	
0	0.64	0.30	0.41	559
1	0.47	0.78	0.59	441
accuracy		0.52	1000	
macro avg	0.55	0.54	0.50	1000
weighted avg	0.56	0.52	0.49	1000

```

# Import the required packages
import numpy as np
from sklearn.utils import shuffle
from sklearn import datasets
from sklearn import preprocessing
from sklearn.svm import SVR
from sklearn.metrics import confusion_matrix, classification_report,
mean_squared_error, explained_variance_score
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

# Load the housing dataset and shuffle the data so that you don't bias
your analysis.
data = datasets.fetch_california_housing()
X, y = shuffle(data.data, data.target, random_state=7)
# Use a smaller subset for faster training
X = X[:5000] # Use only first 5000 samples
y = y[:5000]

# Split the dataset into training and testing in an 80/20 format:
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=7)

# Create and train the Support Vector Regressor using a linear kernel.
print("Training SVR model...")
sv_regressor = SVR(kernel='linear', C=1.0, max_iter=1000) # Add max_iter
for faster training
sv_regressor.fit(X_train, y_train)
print("Training completed!")

# Run the regressor on the testing data and predict the output (predicted
labels).
y_test_pred = sv_regressor.predict(X_test)

# Evaluate the performance of the regressor and print the initial metrics.
print("Mean squared error: ", mean_squared_error(y_test, y_test_pred))
print("Explained variance score: ", explained_variance_score(y_test,
y_test_pred))

```

```

# binarize the predicted values & the actual values using threshold of
25.00.
threshold = 2.0
y_pred_label = (y_test_pred > threshold).astype(int)
y_test_label = (y_test > threshold).astype(int)

# Create the confusion matrix using the predicted labels and the actual
labels.
confusion_mat = confusion_matrix(y_test_label, y_pred_label)

# Visualize the confusion matrix.
print("\nConfusion Matrix:")
print(confusion_mat)
plt.imshow(confusion_mat, interpolation='nearest', cmap=plt.cm.gray)
plt.title('Confusion matrix')
plt.colorbar()
ticks = np.arange(2)
plt.xticks(ticks, ticks)
plt.yticks(ticks, ticks)
plt.ylabel('True labels')
plt.xlabel('Predicted labels')
plt.show()

# Print the classification report based on the confusion matrix.
print(classification_report(y_test_label, y_pred_label))
print(f"\nThreshold: 2.0")
print(f"Actual labels - Class 0: {sum(y_test_label == 0)}, Class 1:
{sum(y_test_label == 1)}")
print(f"Predicted labels - Class 0: {sum(y_pred_label == 0)}, Class 1:
{sum(y_pred_label == 1)}")

```

## Explanation and Approach

### 1. Data loading and shuffling

- Uses `sklearn.datasets.fetch\_california\_housing()` to load the California Housing dataset.
- Shuffles the dataset to avoid ordering bias using `sklearn.utils.shuffle`.

### 2. Subsampling for speed

- For faster experimentation, the script reduces the dataset to the first 5,000 samples (this keeps runtime small during local testing).

### 3. Train / Test split

- Splits the data into training and testing sets with an 80/20 ratio using `train\_test\_split`.

### 4. Model training

- Trains a Support Vector Regressor (`sklearn.svm.SVR`) with a linear kernel and C=1.0.
- `max\_iter` is set to 1000 in the script to limit training time when experimenting locally.

### 5. Regression evaluation

- Evaluates the regressor using Mean Squared Error (MSE) and Explained Variance Score (EVS).
- Prints the regression metrics.

### 6. Binarization and classification analysis (post-hoc)

- Binarizes continuous predictions and ground-truth values with a threshold (default `2.0`, meaning \$200k as the scale is in hundreds of thousands).
  - Produces predicted labels and actual labels: 0 = low price ( $\leq$  threshold), 1 = high price ( $>$  threshold).
  - Computes a confusion matrix and prints it.
  - Visualizes the confusion matrix (grayscale heatmap) and prints a classification report (precision / recall / f1-score).